

# class08\_miniproject

## Table of contents

Background . . . . .	1
Principal Component Analysis (PCA) . . . . .	6
Hierarchical clustering . . . . .	12
Combining methods . . . . .	13
Sensitivity . . . . .	16
Prediction . . . . .	17

## Background

In today's class we will be employing all the R techniques for data analysis we have learned thus far - including the machine learning methods of clustering and PCA - to analyze real breast cancer biopsy data.

```
fna.data <- "~/Documents/BIMM143/Class08_WisconsinCancer.csv"  
wisc.df <- read.csv(fna.data, row.names = 1)
```

wee peak at the data

```
head(wisc.df, 4)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017

```

84300903      0.10960      0.15990      0.1974      0.12790
84348301      0.14250      0.28390      0.2414      0.10520
               symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
842302        0.2419       0.07871     1.0950     0.9053      8.589
842517        0.1812       0.05667     0.5435     0.7339      3.398
84300903      0.2069       0.05999     0.7456     0.7869      4.585
84348301      0.2597       0.09744     0.4956     1.1560      3.445
               area_se smoothness_se compactness_se concavity_se concave.points_se
842302        153.40      0.006399    0.04904     0.05373     0.01587
842517        74.08       0.005225    0.01308     0.01860     0.01340
84300903      94.03       0.006150    0.04006     0.03832     0.02058
84348301      27.23       0.009110    0.07458     0.05661     0.01867
               symmetry_se fractal_dimension_se radius_worst texture_worst
842302        0.03003     0.006193    25.38       17.33
842517        0.01389     0.003532    24.99       23.41
84300903      0.02250     0.004571    23.57       25.53
84348301      0.05963     0.009208    14.91       26.50
               perimeter_worst area_worst smoothness_worst compactness_worst
842302        184.60      2019.0      0.1622      0.6656
842517        158.80      1956.0      0.1238      0.1866
84300903      152.50      1709.0      0.1444      0.4245
84348301      98.87       567.7      0.2098      0.8663
               concavity_worst concave.points_worst symmetry_worst
842302        0.7119       0.2654      0.4601
842517        0.2416       0.1860      0.2750
84300903      0.4504       0.2430      0.3613
84348301      0.6869       0.2575      0.6638
               fractal_dimension_worst
842302          0.11890
842517          0.08902
84300903        0.08758
84348301        0.17300

```

```
wisc.data <- wisc.df[,-1]
```

Q1. How many observations are in this dataset?

```
nrow(wisc.df)
```

```
[1] 569
```

Q2. How many of the observations have a malignant diagnosis

```
sum(wisc.df$diagnosis == "M")
```

```
[1] 212
```

```
table(wisc.df$diagnosis)
```

```
B      M  
357 212
```

Q3. How many variables/ features in the data are suffixed with \_mean

```
colnames(wisc.df)
```

```
[1] "diagnosis"                  "radius_mean"  
[3] "texture_mean"              "perimeter_mean"  
[5] "area_mean"                 "smoothness_mean"  
[7] "compactness_mean"          "concavity_mean"  
[9] "concave.points_mean"       "symmetry_mean"  
[11] "fractal_dimension_mean"    "radius_se"  
[13] "texture_se"                "perimeter_se"  
[15] "area_se"                   "smoothness_se"  
[17] "compactness_se"            "concavity_se"  
[19] "concave.points_se"         "symmetry_se"  
[21] "fractal_dimension_se"      "radius_worst"  
[23] "texture_worst"             "perimeter_worst"  
[25] "area_worst"                "smoothness_worst"  
[27] "compactness_worst"          "concavity_worst"  
[29] "concave.points_worst"       "symmetry_worst"  
[31] "fractal_dimension_worst"
```

```
length(grep("_mean", colnames(wisc.df)))
```

```
[1] 10
```

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst

6.573234e-02

6.186747e-02

1.806127e-02

We need to remove the `diagnosis` column before we do any further analysis of this dataset - we don't want to pass this to PCA etc. We will save it as a separate vector that we can use later to compare our findings to those of experts.

```
diagnosis <- wisc.df$diagnosis  
wisc.data <- wisc.df[,-1]
```

```
head(wisc.data)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
842302	17.99	10.38	122.80	1001.0	0.11840
842517	20.57	17.77	132.90	1326.0	0.08474
84300903	19.69	21.25	130.00	1203.0	0.10960
84348301	11.42	20.38	77.58	386.1	0.14250
84358402	20.29	14.34	135.10	1297.0	0.10030
843786	12.45	15.70	82.57	477.1	0.12780
	compactness_mean	concavity_mean	concave.points_mean	symmetry_mean	
842302	0.27760	0.3001	0.14710	0.2419	
842517	0.07864	0.0869	0.07017	0.1812	
84300903	0.15990	0.1974	0.12790	0.2069	
84348301	0.28390	0.2414	0.10520	0.2597	
84358402	0.13280	0.1980	0.10430	0.1809	
843786	0.17000	0.1578	0.08089	0.2087	
	fractal_dimension_mean	radius_se	texture_se	perimeter_se	area_se
842302	0.07871	1.0950	0.9053	8.589	153.40
842517	0.05667	0.5435	0.7339	3.398	74.08
84300903	0.05999	0.7456	0.7869	4.585	94.03
84348301	0.09744	0.4956	1.1560	3.445	27.23
84358402	0.05883	0.7572	0.7813	5.438	94.44
843786	0.07613	0.3345	0.8902	2.217	27.19
	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	0.006399	0.04904	0.05373	0.01587	
842517	0.005225	0.01308	0.01860	0.01340	
84300903	0.006150	0.04006	0.03832	0.02058	
84348301	0.009110	0.07458	0.05661	0.01867	
84358402	0.011490	0.02461	0.05688	0.01885	
843786	0.007510	0.03345	0.03672	0.01137	
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	

84300903	0.02250	0.004571	23.57	25.53
84348301	0.05963	0.009208	14.91	26.50
84358402	0.01756	0.005115	22.54	16.67
843786	0.02165	0.005082	15.47	23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.60	2019.0	0.1622	0.6656
842517	158.80	1956.0	0.1238	0.1866
84300903	152.50	1709.0	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663
84358402	152.20	1575.0	0.1374	0.2050
843786	103.40	741.6	0.1791	0.5249
	concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119	0.2654	0.4601	
842517	0.2416	0.1860	0.2750	
84300903	0.4504	0.2430	0.3613	
84348301	0.6869	0.2575	0.6638	
84358402	0.4000	0.1625	0.2364	
843786	0.5355	0.1741	0.3985	
	fractal_dimension_worst			
842302		0.11890		
842517		0.08902		
84300903		0.08758		
84348301		0.17300		
84358402		0.07678		
843786		0.12440		

## Principal Component Analysis (PCA)

The main function in base R is called `prcomp()` we will use the optional argument `scale=TRUE` here as the data columns/features/dimensions are in very different scales in the original data set.

```
wisc.pr <- prcomp(wisc.data, scale = T)

attributes(wisc.pr)

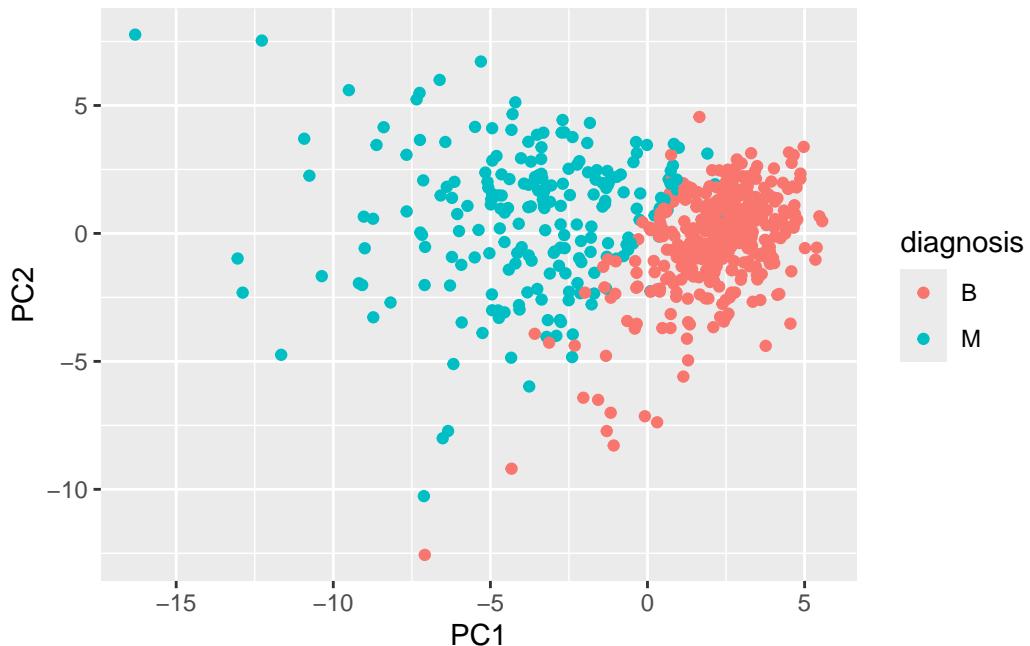
$names
[1] "sdev"      "rotation"   "center"    "scale"     "x"

$class
[1] "prcomp"
```

```

library(ggplot2)
ggplot(wisc.pr$x) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()

```



```

summary(wisc.pr)

```

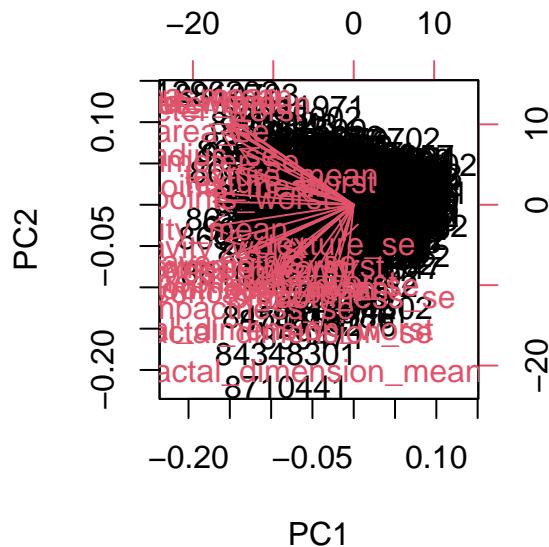
Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987

Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)? 44.27% Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data? Three Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data? Seven

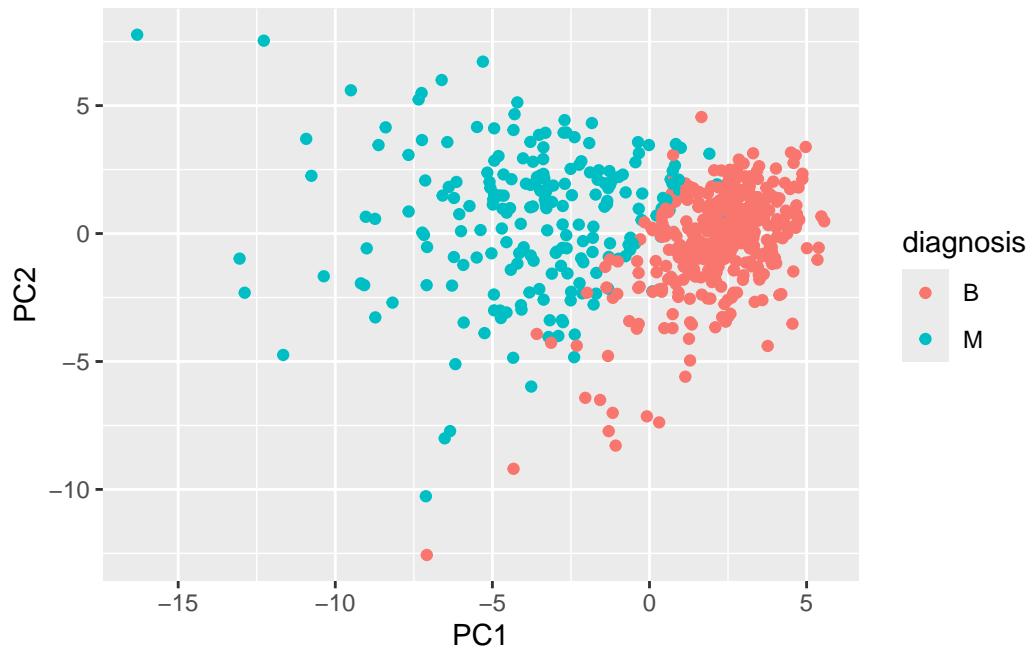
```
biplot(wisc.pr)
```



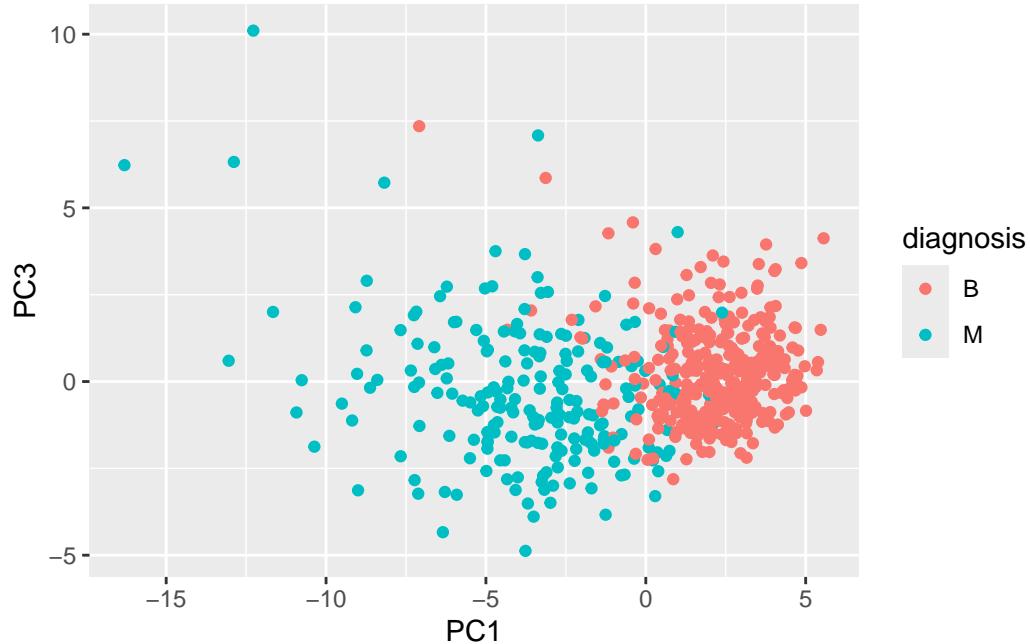
Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

It is very difficult to understand because all the data nad labels are overlapping so nothing can be read/ discerned.

```
library(ggplot2)
ggplot(wisc.pr$x) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



```
library(ggplot2)
ggplot(wisc.pr$x) +
  aes(PC1, PC3, col=diagnosis) +
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

M values are separated from B values by PC1. B and M are also less segregated in the second one.

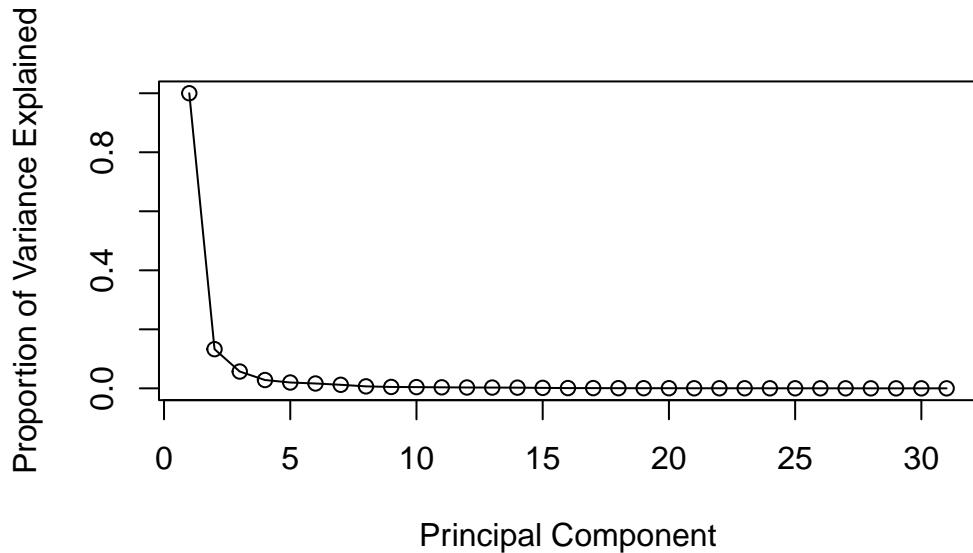
```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

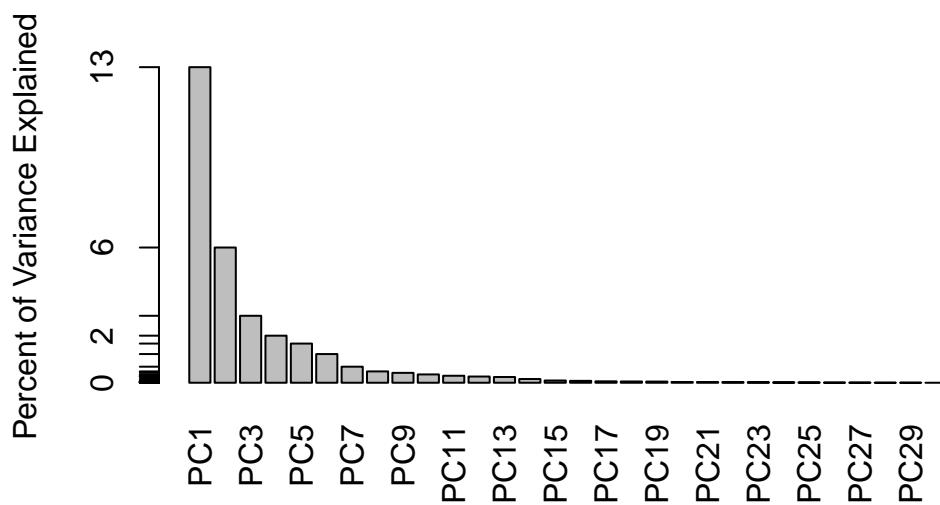
```
pve <- pr.var / 100
head(pve)
```

```
[1] 0.13281608 0.05691355 0.02817949 0.01980640 0.01648731 0.01207357
```

```
plot(c(1, pve), xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



```
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = F)
axis(2, at=pve, labels=round(pve,2)*100)
```



```
wisc.pr$rotation[,1]
```

radius_mean	texture_mean	perimeter_mean
-0.21890244	-0.10372458	-0.22753729
area_mean	smoothness_mean	compactness_mean
-0.22099499	-0.14258969	-0.23928535
concavity_mean	concave.points_mean	symmetry_mean
-0.25840048	-0.26085376	-0.13816696
fractal_dimension_mean	radius_se	texture_se
-0.06436335	-0.20597878	-0.01742803
perimeter_se	area_se	smoothness_se
-0.21132592	-0.20286964	-0.01453145
compactness_se	concavity_se	concave.points_se
-0.17039345	-0.15358979	-0.18341740
symmetry_se	fractal_dimension_se	radius_worst
-0.04249842	-0.10256832	-0.22799663
texture_worst	perimeter_worst	area_worst
-0.10446933	-0.23663968	-0.22487053
smoothness_worst	compactness_worst	concavity_worst
-0.12795256	-0.21009588	-0.22876753
concave.points_worst	symmetry_worst	fractal_dimension_worst
-0.25088597	-0.12290456	-0.13178394

Q9. For the first principal component, what is the component of the loading vector (i.e. wisc.pr\$rotation[1]) for the feature concave.points\_mean? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

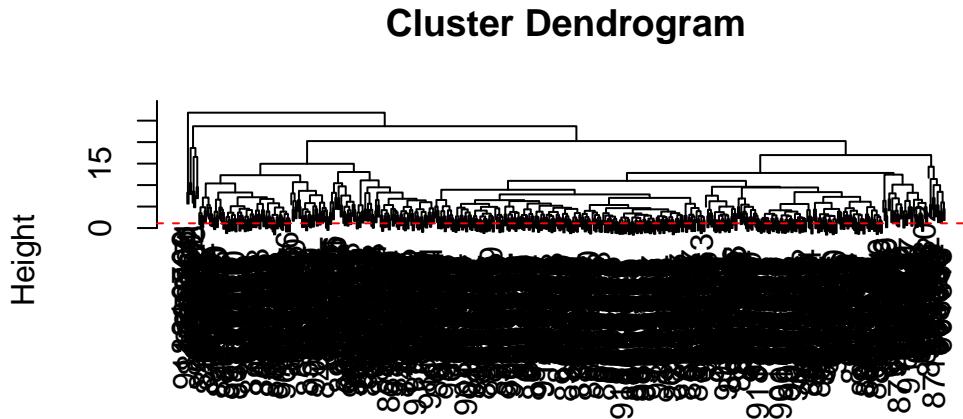
-0.26085376. No there are not.

## Hierarchical clustering

```
data.scaled <- scale(wisc.data)
data.list <- dist(data.scaled)
wisc.hclust <- hclust(data.list, "complete")
```

Q10. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h = wisc.hclust$height[4], col="red", lty=2)
```



```
data.list
hclust (*, "complete")
```

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=4)
table(wisc.hclust.clusters, diagnosis)
```

diagnosis		
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

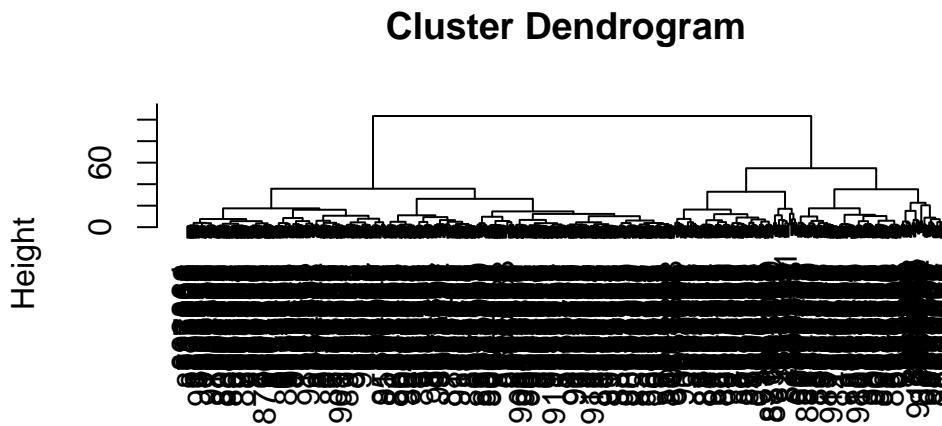
Q12. Which method gives your favorite results for the same data.dist dataset?  
Explain your reasoning

ward.D2 because it puts the B and M values in different clusters to the greatest extent.

### Combining methods

The idea here is that I can take my new variables (i.e. the PCs) that are better descriptors of the data-set than the original features (i.e. the 30 columns in `wisc.data`) and use these as a basis for clustering.

```
pc.dist <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(pc.dist, method = "ward.D2")
plot(wisc.pr.hclust)
```



pc.dist  
hclust (\*, "ward.D2")

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1 2
203 366
```

```
table(diagnosis)
```

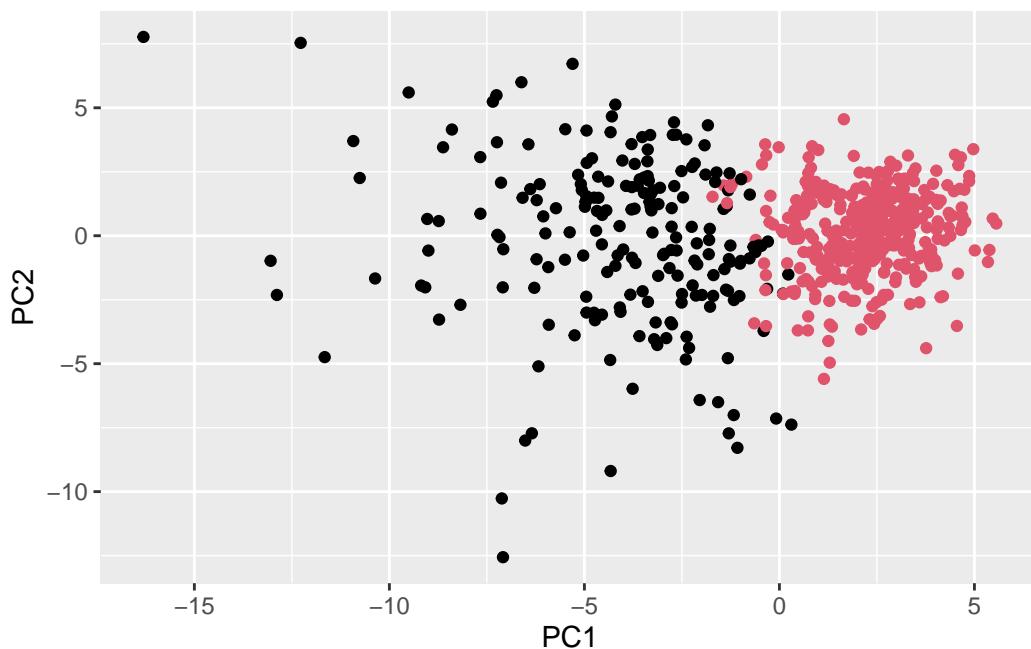
```
diagnosis
  B   M
357 212
```

I can now run `table()` with both my clustering `grps` and the expert `diagnosis`

```
table(grps, diagnosis)
```

	diagnosis	
grps	B	M
1	24	179
2	333	33

```
ggplot(wisc.pr$x) +  
  aes(PC1, PC2) +  
  geom_point(col=grps)
```



```
hclust_dist <- dist(wisc.pr$x[, 1:7])  
wisc.pr.hclust <- hclust(hclust_dist, method="ward.D2")  
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
```

Q13. How well does the newly created hclust model with two clusters separate out the two “M” and “B” diagnoses?

```
table(wisc.pr.hclust.clusters, diagnosis)
```

```

diagnosis
wisc.pr.hclust.clusters   B   M
                           1 28 188
                           2 329 24

```

Fairly well, though there are still 52 misassigned diagnoses.

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.hclust.clusters and wisc.pr.hclust.clusters) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters, diagnosis)
```

```

diagnosis
wisc.hclust.clusters   B   M
                           1 12 165
                           2  2  5
                           3 343 40
                           4  0  2

```

Again, pretty well, the Benign diagnoses are more accurate, but the malignant diagnoses are less accurate.

## Sensitivity

Our cluster “1” has 24 benign and 179 malignant, cluster “2” has 333 benign and 33 malignant.

179 TP (true positives) 24 FP 333 TN (true negative) 33 FN

Sensitivity:  $\frac{TP}{TP+FN}$

```
179/(179+33)
```

```
[1] 0.8443396
```

Sensitivity of 1 is perfect.

Specificity:  $\frac{TN}{TN+FP}$

```
333/(333+24)
```

```
[1] 0.9327731
```

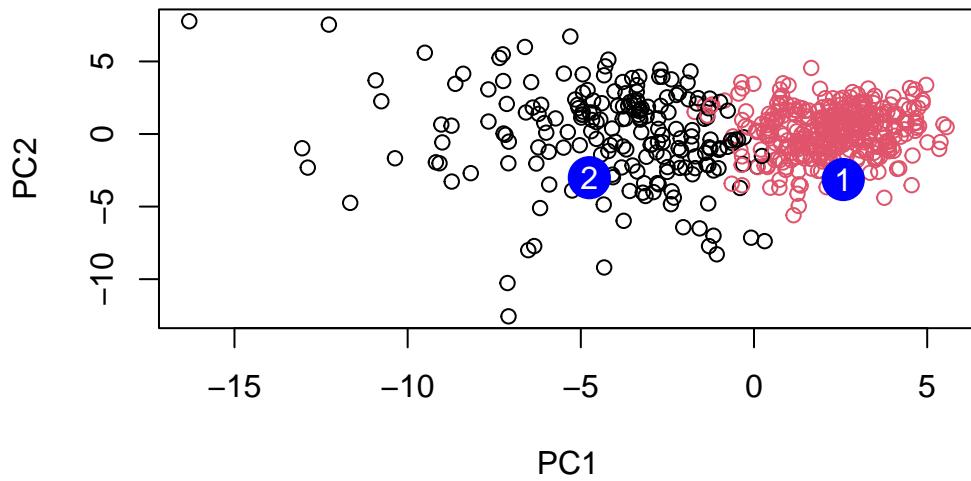
## Prediction

We can use our PCA model for prediction of new unseen data.

```
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q16. Which of these new patients should we prioritize for follow up based on your results?

Group two since they have the malignant diagnoses.